

Docker Tutorial

1. Containers, Docker, Docker Compose
2. Docker Registry
3. Kubernetes
 - a. Pods
 - b. Workloads
 - c. Services
 - d. Updates
 - e. Storage
 - f. App Settings
 - g. Observability
 - h. Scaling

<https://youtu.be/kTp5xUcalw>

<https://docs.docker.com/engine/install/ubuntu/>

<https://landscape.cncf.io/>

Why containers ?

- Move faster by deploying smaller units
- Use fewer resources
- Fit more into the same host
- Faster automation
- Portability
- Isolation

Container Registry

Orchestrator

- Manage
 - Infrastructure
 - Containers
 - Deployment
 - Scaling
 - Failover
 - Health monitoring
 - App upgrades, Zero-Downtime deployments
- Install your own
 - Kubernetes, Swarm, Service Fabric
- Orchestrators as a service
 - Azure Kubernetes Service, Service Fabric

What is Docker?

- An open source container runtime
- Mac, Windows & Linux support
- Command line tool
- Dockerfile file format for building container images

Docker commands

- i. `docker pull [imageName]` - Pull an image from a registry
- j. `docker run [imageName]` - Run containers
- k. `docker run -d [imageName]` - detached mode
- l. `docker start [containerName]` - start stopped containers
- m. `docker ps` - list running containers
- n. `docker ps -a` - list running and stopped containers
- o. `docker stop [containerName]` - stop containers
- p. `docker kill [containerName]` - kill containers
- q. `docker image inspect [imageName]` - get image info
- r. `docker rm [containerName]`
- s. `docker run -it nginx -- /bin/bash`
- t. `docker run -it -- microsoft/powershell:nanoserver pwsh.exe` → attach powershell
- u. `docker container exec -it [containerName] --bash` → Attach to a running container
- v. `docker rm $(docker ps -a -q)` → Removes all stopped containers
- w. `docker images` → Lists images
- x. `docker rm [imageName]`
- y. `docker system prune -a` → Removes all images not in use by any containers

imageName - name of the image on the repo

`docker run --publish 80:80 --name webserver nginxdo`

Build containers

- `docker build -t [name:tag] .` → Builds an image using a Dockerfile located in the same folder
- `docker build -t [name:tag] folder` → Builds an image using a Dockerfile located in a different folder
- `docker tag [imageName] [name:tag]` → Tag an existing image
- **Example: Static HTML SITE**
 - FROM nginx:alpine
 - COPY . /usr/share/nginx/html
 - `docker build -t webserver-image:v1 .` → build
 - `docker run -d -p 8080:80 webserver-image:v1` → run
 - Curl localhost:8000 → display

- **Example : Dockerfile - Node site**
 - FROM alpine
 - RUN apk add -update nodejs nodejs-npm
 - COPY . /src
 - WORKDIR /src
 - RUN npm install
 - EXPOSE 8080
 - ENTRYPOINT ["node", "./app.js"]

Docker tagging

- docker tag → create a target image
 - name:tag
 - myimage:v1
 - repository/name:tag
 - myacr.azurecr.io/myimage:v1

Docker tutorial by Mosh

1. What is Docker

A platform for building, running and shipping applications

2. What is a Container

- allows running multiple apps in isolation
- Are lightweight
- Use OS of the host
- Start quickly
- Need less hardware resources

3. Docker Architecture

4. Docker Commands

- a. Build an image
 - i. docker build -t hello-docker, `hello-docker` is a tag name
- b. List images
 - i. docker image ls
 - ii. docker images

Docker tutorial by **Academind**

1. Module 1 & 2

a. Docker intro

- i. Image vs Container
 - 1. Images - templates/ blueprints for containers
 - 2. Containers - the running unit of software
- ii. Images
 - 1. Images are layer based and each layer is cached. Each instruction corresponds to a layer.
 - 2. Images are read only

b. Docker commands

- i. Build an image
 - 1. `docker build [path to the Dockerfile]`
 - 2. `docker build -t [nameOfTheImage] .`
- ii. Managing Images & Containers
 - 1. Running the containers
 - a. `docker run -p 3000:80 [imageName/Id]`
 - b. `docker run -p 3000:80 -d [imageName/Id]` → running in a detached mode
 - c. `docker run -p 3000:80 -d --rm [imageName/Id]` → running in a detached mode and removes the container when exited
 - 2. List containers
 - a. `docker ps`
 - b. `docker ps -a` → list stopped containers too
 - 3. Restart a stopped container
 - a. `docker start [containerName]`
 - b. `docker start -a [containerName]` → in attached mode
 - 4. Attach to a container
 - a. `docker attach [containerName/id]`
 - 5. View the logs
 - a. `docker logs` → view all the logs
 - b. `docker logs -f` → similar to `tail -f`, you can see the logs in realtime
 - 6. Attach in interactive mode
 - a. `docker run -it [imageName/id]`
 - b. `docker start -a -i [containerName]`
 - 7. Deleting images & Containers
 - a. `docker rm [containerName] [containerName]`
 - b. `docker rmi [imageName]`
 - c. `docker image prune` → remove unused containers

8. Inspecting the image
 - a. `docker image inspect [imageName/id]`
9. Copying files into and from a container
 - a. From local machine to container
 - i. `docker cp dummy/. jolly_carson:/test → docker cp [sourceFolder/filePath] [containerName]:[pathInsideTheContainer]`
 - b. From container to a local machine
 - i. `docker cp jolly_carson:/test/test.txt dummy/ → docker cp [containername]:[pathInsideTheContainer] [destinationPath]`
10. Naming containers and images
 - a. Naming a container
 - i. `docker run -p 3000:80 -d --rm --name [anyNameForTheContainer] [imageName/Id]`
 - b. Naming/tagging an image
 - i. `docker build -t [anyName]:[tag/version] .`
 - ii. Renaming
 1. `docker tag node-hello-world:latest mahatoniteesh/node-hello-world`
11. Pushing and pulling images to/from docker hub
 - a. `docker login → login to the hub`
 - b. `docker push [imageName] → image name should be same as the repository name create online e.g. mahatoniteesh/node-app`
 - c. `docker pull [repositoryName]:[version]`

2. Module 3

a. Volumes

- i. Are folder that reside on the host machine
- ii. These are used by container to store files
- iii. When the container is killed/stopped, the data persists
- iv. Types:
 1. Anonymous Volume
 - a. `docker run -v /app/data...`
 2. Name volume
 - a. `docker run -v data:/app/data..`
 3. Bind Mount
 - a. `docker run -v /path/to/code:/app/code...`

b. Named Volumes

- i. Creating name volume
 1. `docker run --name feedback-app -d -p 3000:80 -v feedback:/app/feedback feedback:volume → -v flag is used to`

attach a name volume to the container which is persisted when the container gets terminated

2. `-v [volumeName]:[pathInsideTheContainerToLink]`
- ii. Listing the volumes
 1. `docker volume ls`
- iii. Removing the volumes
 1. `docker volume rm [volumeName]`
 2. `docker volume prune`
- c. Bind Mounts (Code sharing)
 - i. Creating
 1. `docker volume create [volumeName]`
 2. `-v $(pwd):/app` → create while container run command
 - ii. Running container with bind mount
 1. `docker run --name feedback-app -d -p 3000:80 -v feedback:/app/feedback -v $(pwd):/app -v /app/node_modules feedback:volume`
 - iii. Read only volume
 1. `-v feedback:/app/feedback -v $(pwd):/app:ro` → add `ro`
- d. ENV variables
 - i. Provide env
 1. In docker file
`ENV PORT 80`
`EXPOSE $PORT`
 2. `-e PORT=3000`
 - ii. Provide env file
 1. `--env-file .env`
 - iii. While running container
 1. `docker run --name feedback-app -e PORT=8000 -p 3000:8000 -d -v $(pwd):/app -v /app/node_modules -v feedback:/app/feedback --rm feedback-app:env`
- e. ARG variables
 - i. In docker file
 1. `ARG DEFAULT_PORT=80`
`ENV PORT $DEFAULT_PORT`
 2. `docker build --build-arg DEFAULT_PORT=8000`

3. Module 4 - Networking

- a. Accessing localhost
 - i. `host.docker.internal` → localhost on the host machine
- b. Container to Container
 - i. Using container ip address → inefficient
 - ii. Using docker network
 1. `docker network create [networkName]`
 - iii. Adding container to a network

1. `docker run --network [networkName]` → you can connect to other containers using the containerName [http://\[containerName\]](http://[containerName])

4. Module 5 - Multi Container apps

5. Module 6 - Docker Compose

a. Installing

- i. On macOS and Windows, you should already have Docker Compose installed - it's set up together with Docker there.
- ii. On Linux machines, you need to install it separately.
- iii. These steps should get you there:
- iv. 1. `sudo curl -L "https://github.com/docker/compose/releases/download/1.27.4/docker-compose-$(uname -s)-$(uname -m)" -o /usr/local/bin/docker-compose`
- v. 2. `sudo chmod +x /usr/local/bin/docker-compose`
- vi. 3. `sudo ln -s /usr/local/bin/docker-compose /usr/bin/docker-compose`
- vii. 4. to verify: `docker-compose --version`
- viii. Also see: <https://docs.docker.com/compose/install/>

b. Writing docker compose file

version: "3.8"

```
services:
  mongodb:
    image: "mongo"
    volumes:
      - ./mongo-data:/data/db
    # environment:
    #   - MONGO_INITDB_ROOT_USERNAME=$MONGO_USER
    #   - MONGO_INITDB_ROOT_PASSWORD=$MONGO_PASS
    env_file:
      - ./env/mongo.env
    # networks:
    #   - network
    # container_name: mongodb
backend:
  # build: path
  # build:
```

```

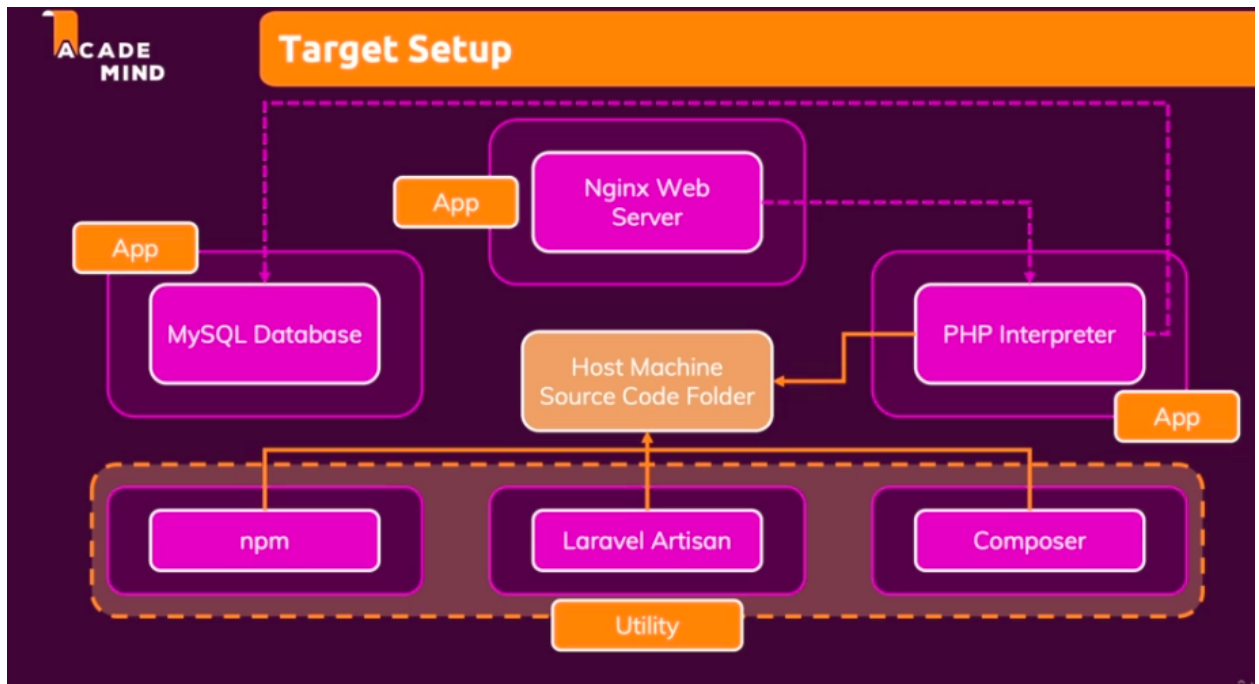
#   context: path
#   dockerfile: Dockerfile
#   args:
#     - name=value
ports:
  - "80:3000"
volumes:
  - m4backend-logs:/app/logs
  - ./backend:app
  - /app/node_modules
env_file:
  - ./env/backend.env
depends_on:
  - mongodb
frontend:
  build: ./frontend
  ports:
    - '3000:3000'
  volumes:
    - ./frontend/src:app/src
  stdin_open: true
  tty: true
  depends_on:
    - backend
volumes:
  m4backend:

```

6. Module 7 - utility containers

- a. `docker exec -it <container-name> npm init` → to execute a command inside the container
- b. `docker run -it <image-name> <command-to-run>` → `docker run -it node npm init`
- c. You can specify entrypoint to restrict the command
 - i. `ENTRYPOINT ["npm"]`
 - ii. `docker run -t node init` → which will be then resolved to `npm init`
 - iii. For compose file
 1. `docker compose run <container-name/service-name> init`

7. Module 8 – Laravel php setup



commands :

- Docker-compose up -d --build <service-name> → force build containers

8. Module 9 – Deploying containers

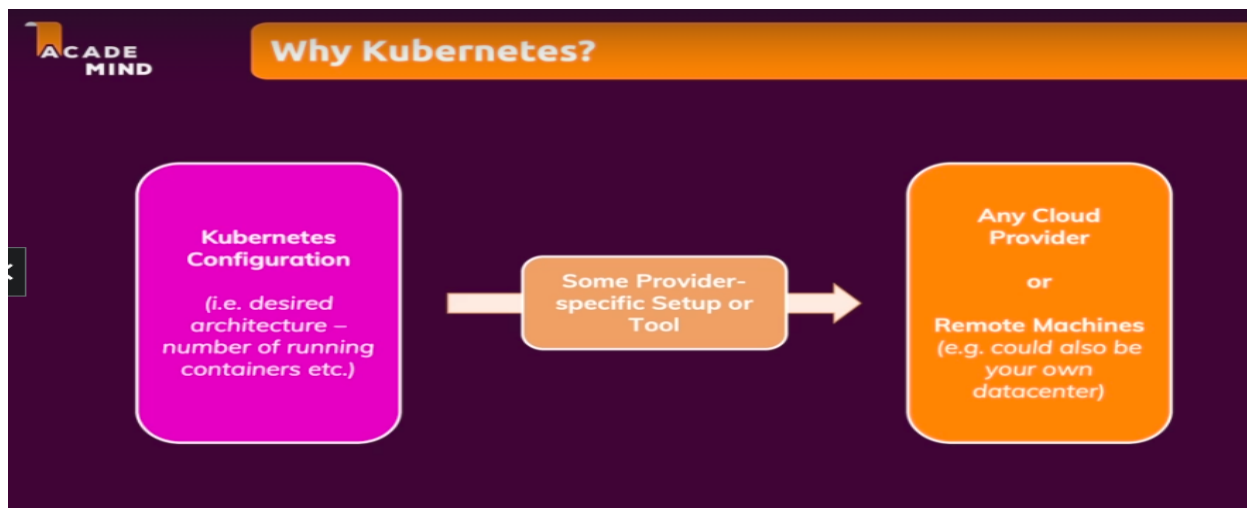
Check cloud formation

Cloudwatch

Kubernetes

1. Introduction

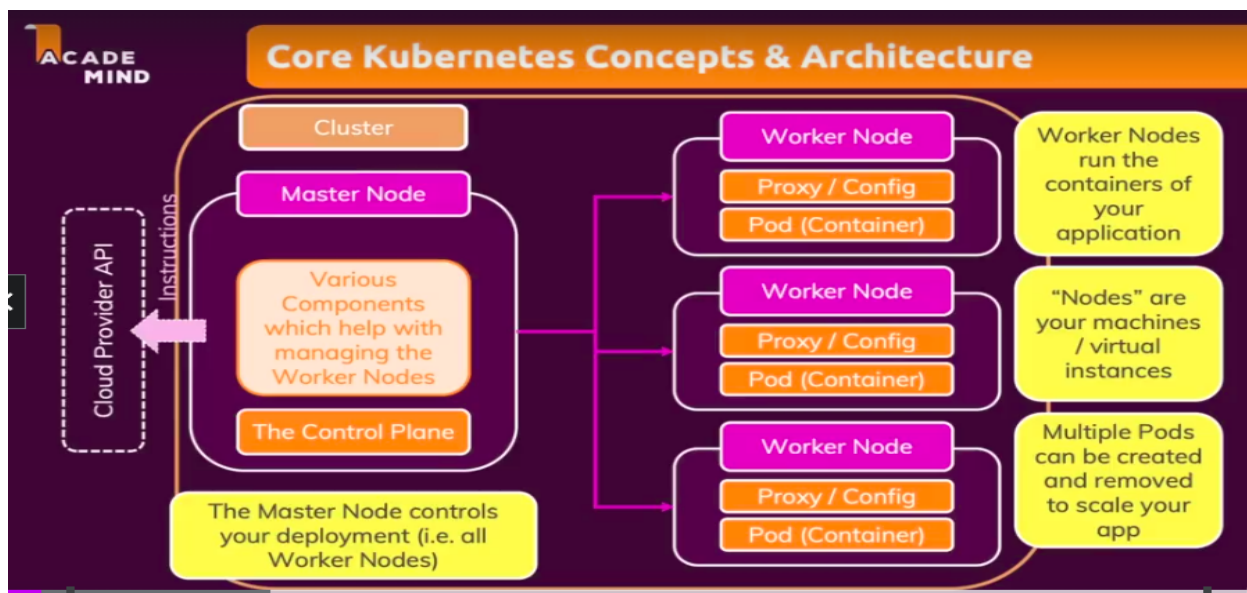
- a. Why Kubernetes ?



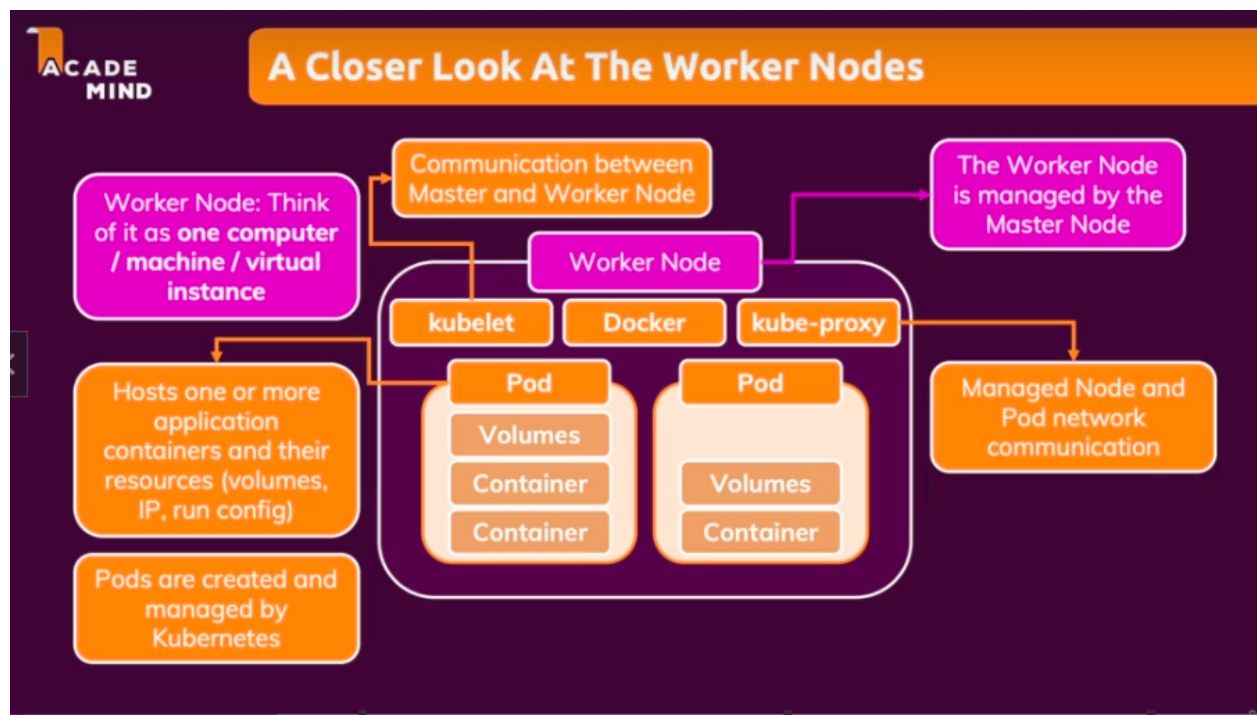
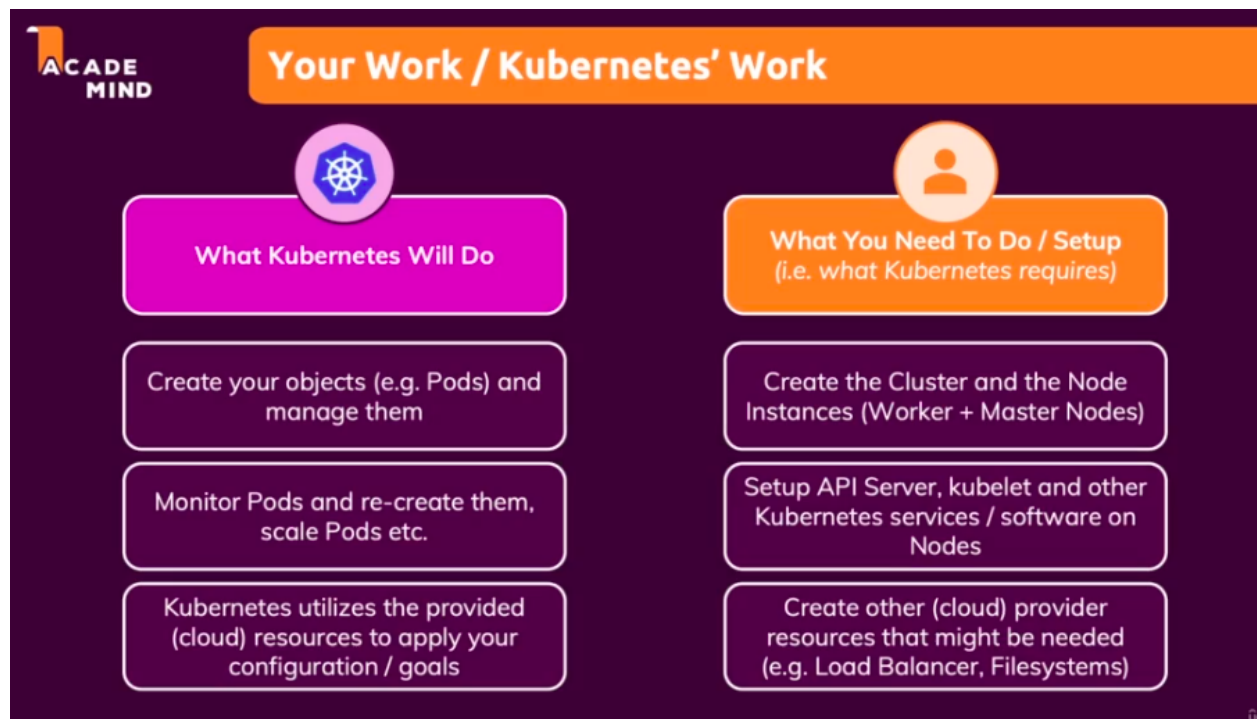
b. Concepts

- i. Pod → Smallest possible unit
 1. Contains 1 or more containers
- ii. Worker node → A virtual machine eg. ec2
 1. On which 1 or more pods run
- iii. Proxy/config
 1. Used for networking.
- iv. Master Node → Another machine
 1. Controls worker node

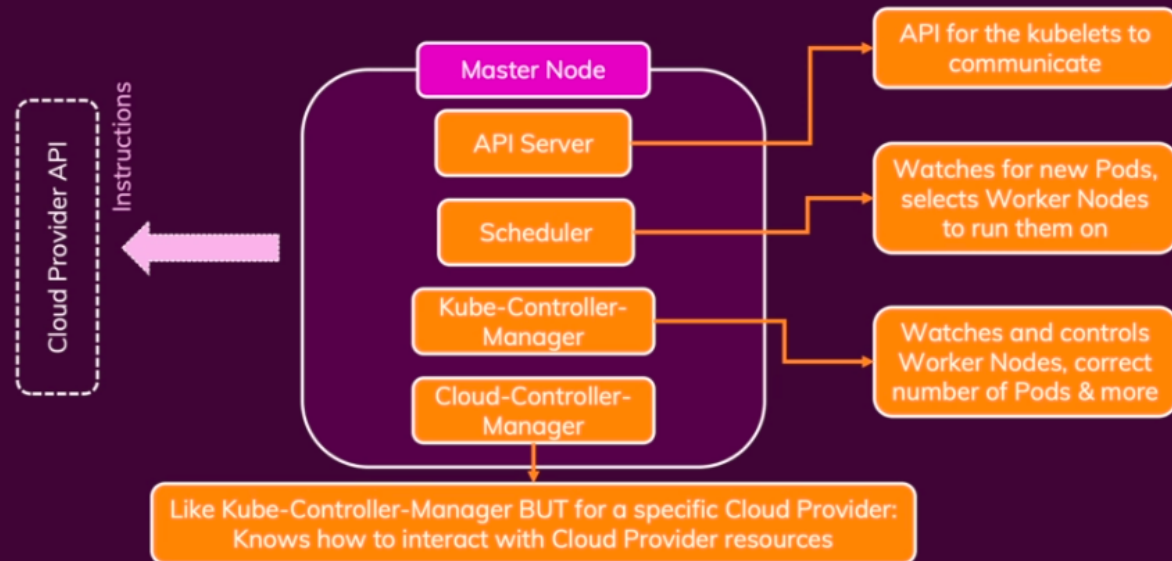
All the above together forms the cluster



Your Work / Kubernetes Work



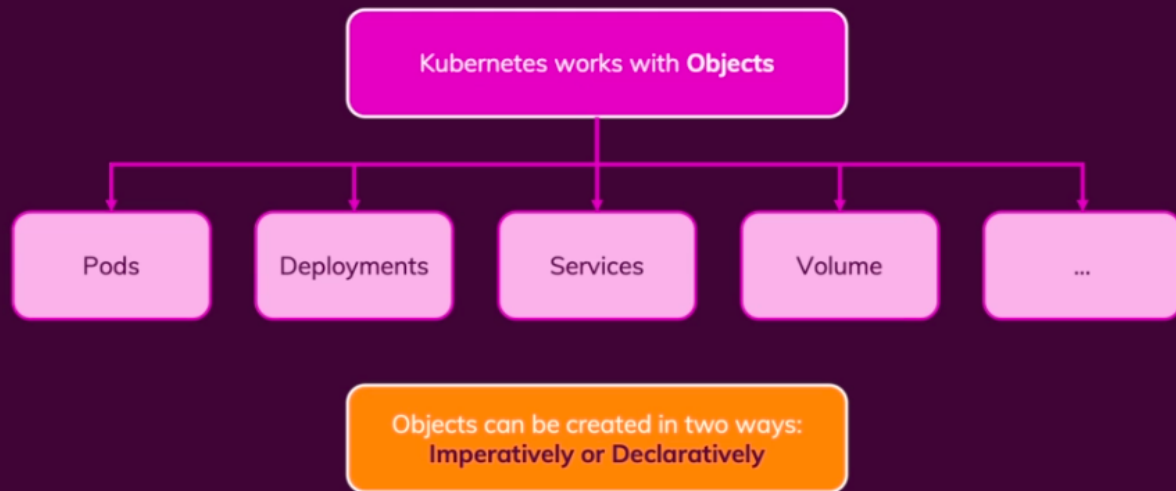
Core Kubernetes Concepts & Architecture



Core Components

Cluster	A set of Node machines which are running the Containerized Application (Worker Nodes) or control other Nodes (Master Node)
Nodes	Physical or virtual machine with a certain hardware capacity which hosts one or multiple Pods and communicates with the Cluster
Master Node	Cluster Control Plane , managing the Pods across Worker Nodes
Worker Node	Hosts Pods , running App Containers (+ resources)
Pods	Pods hold the actual running App Containers + their required resources (e.g. volumes).
Containers	Normal (Docker) Containers
Services	A logical set (group) of Pods with a unique, Pod- and Container-independent IP address

Understanding Kubernetes Objects

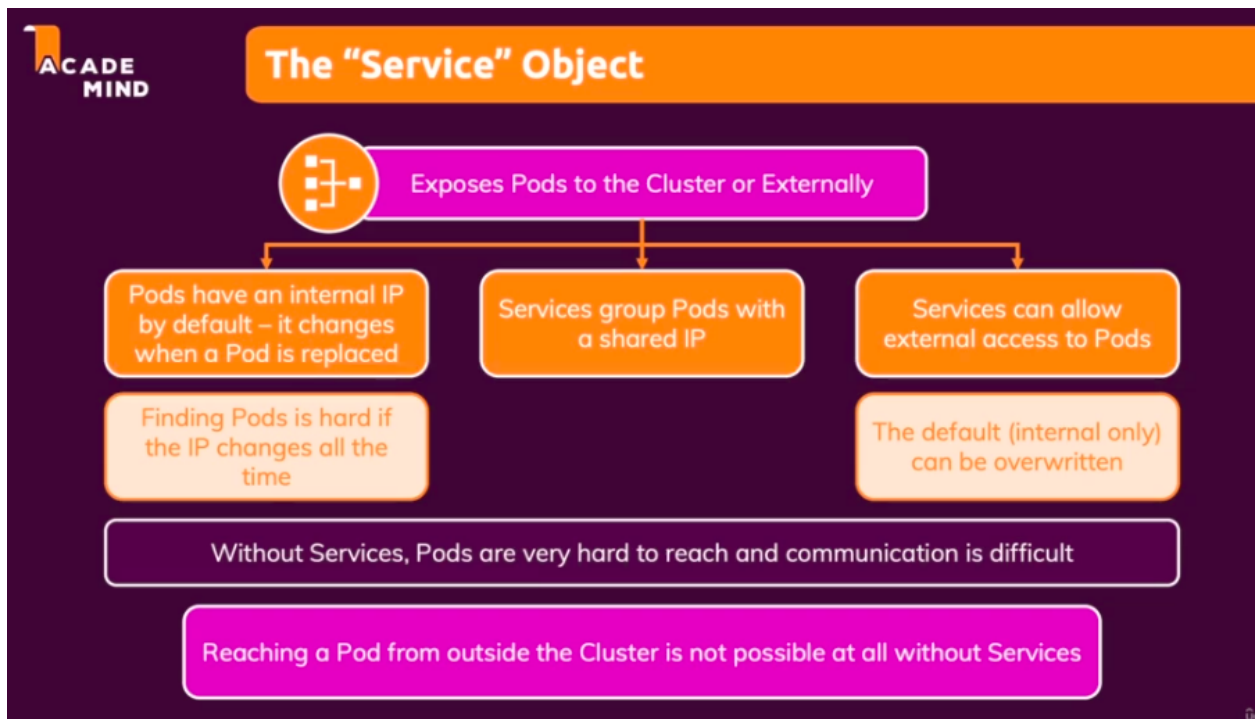


2. Commands

a. Deployment Object

- i. Get deployments
 1. `kubectl get deployments`
- ii. Get pods
 1. `kubectl get pods`
- iii. Delete deployments
 1. `kubectl delete deployment <app-name>`
- iv. Create deployment
 1. `kubectl create deployment first-app --image=<docker-hub-repo-name>`
- v. Web dashboard
 1. minikube dashboard

- b. Service Object → Expose pods to the cluster or Externally.



- i. Creating a service
 1. `kubectl service create`
- ii. Expose a port
 1. `kubectl expose deployment <app-name> --type=<type-of-expose> --port=8080`
 - a. Expose types can be
 - i. ClusterIP → will give an ip that is reachable only inside the cluster
 - ii. NodePort → will give an ip of the the worker node, now the app will be accessible to the outside world
 - iii. LoadBalancer → this will assign a public ip and will also distribute the load
 2. Locally assigning ip to a loadbalancer service
 - a. `minicube service <app-name>`
- iii. List services
 1. `kubectl get services`
- iv. Scaling - create replicas
 1. `kubectl scale deployment/first-app --replicas=3` → will create 3 replica pods of deployment/first-app
 2. `kubectl scale deployment/first-app --replicas=1` → will downgrade to 1 replica pod of deployment/first-app

c. Updating deployments

- i. updating an image
 1. `kubectl set image deployment/first-app kub-first-app=mahatoniteesh/kub-first-app`
- ii. Get rollout status/update status
 1. `kubectl rollout status deployment/first-app`
- iii. Get rollout history of a deployment
 1. `kubectl rollout history deployment/first-app`
- iv. Rolling back to a particular revision
 1. `kubectl rollout undo deployment/app --to-revision=2` → rolling back to the second revision

d. Config files

i. Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: second-app-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app : second-app
      tier : backend
  template:
    metadata:
      labels:
        app: second-app
        tier: backend
    spec:
      containers:
        - name: second-node
          image: mahatoniteesh/kube-first-app
```

ii. Service

```
apiVersion: v1
kind: Service
metadata:
  name: backend
spec:
  selector:
    app: second-app
```

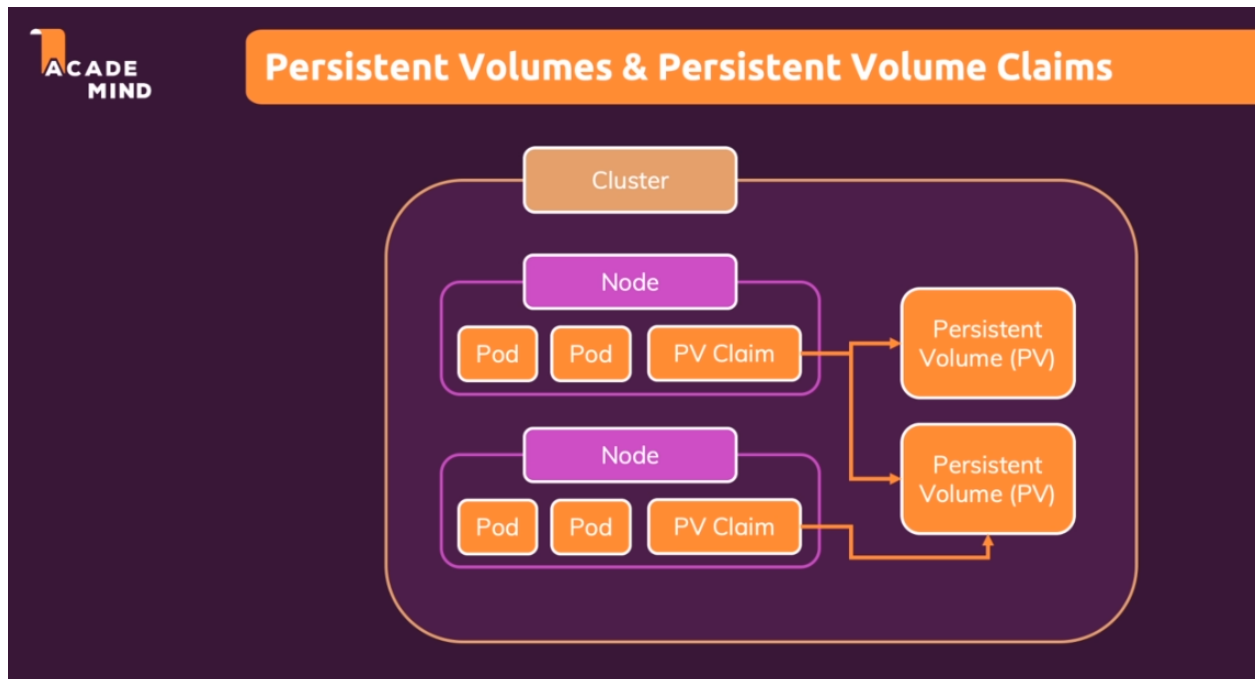
```
ports:
  - protocol : 'TCP'
    port: 80
    targetPort: 8080
type: LoadBalancer
```

iii. Commands

1. Starting from a file
 - a. `kubectl apply -f=deployment.yaml`
2. Delete the resources
 - a. `kubectl delete -f=deployment.yaml -f=... -f=..`
3. Delete by label
 - a. `kubectl delete deployments,services -l group=example`

e. Volume

- i. `emptyDir`
- ii. `hostPath`
- iii. CSI
- iv. Persistent Volume
 1. Get
 - a. `kubectl get pv`
 2. Get storage class
 - a. `kubectl get sc`
 3. Get storage claims
 - a. `kubectl get pvc`



Persistent Volume

```
apiVersion: v1
```



```
kind: PersistentVolume
metadata:
  name: story-pv-volume
spec:
  storageClassName: manual
  capacity:
    storage: 1Gi
  volumeMode: Filesystem
  storageClassName: standard
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/data"
    type: DirectoryOrCreate
```

Persistent volume claim

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: story-pvc
spec:
  volumeName: story-pv-volume
  storageClassName: standard
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
```

Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: story-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: story
  template:
```

```

metadata:
  labels:
    app: story
spec:
  containers:
    - name: story
      image: mahatoniteesh/kub-volumes-demo
      volumeMounts:
        - mountPath: /app/story
          name: story-volume
      env:
        - name: STORY_FOLDER
          value: 'story'
  volumes:
    # - name: story-volume
    #   emptyDir: {}
    # - name: story-volume
    #   hostPath:
    #     path: /data
    #     type: DirectoryOrCreate
    - name: story-volume
      persistentVolumeClaim:
        claimName: story-pvc

```

v. Environment variables and configmap

`spec:`

```

containers:
  - name: story
    image: mahatoniteesh/kub-volumes-demo
    volumeMounts:
      - mountPath: /app/story
        name: story-volume
    env:
      - name: STORY_FOLDER
        value: 'story'
    env:
      - name: STORY_FOLDER
        valueFrom:
          configMapKeyRef:
            name: data-store-env

```

```
key: folder
```

vi. Config maps

```
apiVersion: v1
```

```
kind: ConfigMap
```

```
metadata:
```

```
  name: data-store-env
```

```
data:
```

```
  folder: 'story'
```

f. Networking

- i. To communicate with a container within a pod
 1. Use `localhost`
- ii. To communicate between pods
 1. Create a service for each pod
 2. Connect
 - a. Method 1
 - i. User env var generated by kubernetes
 1. It will be in the form, suppose your service name is **auth-service** then the env variable will be **AUTH_SERVICE_SERVICE_HOST**. So depending on your service name the first part will change but **_SERVICE_HOST** will remain the same.
 - ii. Using internal cluster DNS by kubernetes
 1. They are in the form **<service_name>.<namespace>**
 2. So if the service name is **auth-service** and the namespace is **default** then the dns will be **auth-service.default**
 3. You can get the namespace name by **kubectl get namespaces**

g. Deployment aws eks